

```
In [1]: import pandas as pd
import numpy as np
```

```
In [4]: df = pd.read_csv("C:/Users/pravi/Desktop/Python_April/Automobile_data.csv" )
```

```
In [8]: df.head(5)
```

Out[8]:

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4

5 rows × 26 columns



```
In [11]: df.tail()
```

Out[11]:

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base
200	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1
201	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1
202	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1
203	-1	95	volvo	diesel	turbo	four	sedan	rwd	front	109.1
204	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1

5 rows × 26 columns



```
In [14]: orig_col_names = df.columns
print(orig_col_names)
```

```
Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
      'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
      'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
      'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
      'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
      'highway-mpg', 'price'],
      dtype='object')
```

```
In [17]: # Column headers and mutability
df_new = df.copy()
df_new.columns = ['Symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
                  'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
                  'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
                  'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
                  'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
                  'highway-mpg', 'price']
df_new.head()
```

Out[17]:

	Symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4

5 rows × 26 columns



```
In [ ]: df = df_new.copy()
```

In [18]: `df.head()`

Out[18]:

	Symbols	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4

5 rows × 26 columns



In [20]: `df_test = df[["make", "Symbols"]]
df_test["make"] = 0
df_test.head()`

C:\Users\pravi\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[20]:

	make	Symbols
0	0	3
1	0	3
2	0	1
3	0	2
4	0	2

In [21]: `df.head()`

Out[21]:

	Symbols	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4

5 rows × 26 columns



In [22]: `# export dataframe
df.to_csv("auto.csv", sep = ',')`

In []: `#df["Symbol"].to_list()`

In [23]: `df.iloc[1:4, 4:7]`

Out[23]:

	aspiration	num-of-doors	body-style
1	std	two	convertible
2	std	two	hatchback
3	std	four	sedan

In [26]: `df['Symbols']`

Out[26]:

```

0      3
1      3
2      1
3      2
4      2
...
200    -1
201    -1
202    -1
203    -1
204    -1
Name: Symbols, Length: 205, dtype: int64
```

In [27]: *# Check/Analyze the data*
df.dtypes

Out[27]:

Symbols	int64
normalized-losses	object
make	object
fuel-type	object
aspiration	object
num-of-doors	object
body-style	object
drive-wheels	object
engine-location	object
wheel-base	float64
length	float64
width	float64
height	float64
curb-weight	int64
engine-type	object
num-of-cylinders	object
engine-size	int64
fuel-system	object
bore	object
stroke	object
compression-ratio	float64
horsepower	object
peak-rpm	object
city-mpg	int64
highway-mpg	int64
price	object
dtype:	object

In [28]: df.describe()

Out[28]:

	Symbols	wheel-base	length	width	height	curb-weight	engine-size	co
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	
mean	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317	
std	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693	
min	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	
25%	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000	
50%	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	
75%	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000	
max	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000	

```
In [29]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Symbols              205 non-null    int64
1   normalized-losses    205 non-null    object
2   make                 205 non-null    object
3   fuel-type            205 non-null    object
4   aspiration            205 non-null    object
5   num-of-doors         205 non-null    object
6   body-style           205 non-null    object
7   drive-wheels         205 non-null    object
8   engine-location      205 non-null    object
9   wheel-base          205 non-null    float64
10  length               205 non-null    float64
11  width                205 non-null    float64
12  height               205 non-null    float64
13  curb-weight          205 non-null    int64
14  engine-type          205 non-null    object
15  num-of-cylinders     205 non-null    object
16  engine-size          205 non-null    int64
17  fuel-system          205 non-null    object
18  bore                 205 non-null    object
19  stroke               205 non-null    object
20  compression-ratio    205 non-null    float64
21  horsepower           205 non-null    object
22  peak-rpm             205 non-null    object
23  city-mpg             205 non-null    int64
24  highway-mpg          205 non-null    int64
25  price                205 non-null    object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB
```

```
In [ ]: x = list[1,2]
```

```
In [33]: # Accessing columns
x = np.array(df["body-style"])
```

```
In [35]: # Missing value

df.replace('?', np.nan, inplace=True)
df.head()
```

Out[35]:

	Symbols	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base
0	3	NaN	alfa- romero	gas	std	two	convertible	rwd	front	88.6
1	3	NaN	alfa- romero	gas	std	two	convertible	rwd	front	88.6
2	1	NaN	alfa- romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4

5 rows × 26 columns



```
In [36]: df.shape
```

Out[36]: (205, 26)

```
In [37]: print(df.shape)
df.dropna(subset = ["price"], axis = 0)
print(df.shape)
df.dropna(subset = ["price"], axis = 0, inplace = True)
print(df.shape)
```

(205, 26)

(205, 26)

(201, 26)

```
In [49]: df_null = df.isnull()
df_null.head(5)
```

Out[49]:

	Symbols	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base	...	e
0	False	False	False	False	False	False	False	False	False	False	...	
1	False	False	False	False	False	False	False	False	False	False	...	
2	False	False	False	False	False	False	False	False	False	False	...	
3	False	False	False	False	False	False	False	False	False	False	...	
4	False	False	False	False	False	False	False	False	False	False	...	

5 rows × 26 columns



```
In [51]: # Find columns with missing data and number of missing values
for col in df_null.columns:
    # print(col)
    if True in df_null[col].unique():
        print("in loop", col)
        print(df_null[col].value_counts())
```

```
in loop num-of-doors
False    199
True      2
Name: num-of-doors, dtype: int64
in loop bore
False    197
True      4
Name: bore, dtype: int64
in loop stroke
False    197
True      4
Name: stroke, dtype: int64
in loop horsepower
False    199
True      2
Name: horsepower, dtype: int64
in loop peak-rpm
False    199
True      2
Name: peak-rpm, dtype: int64
```

```
In [44]: # Unique values in a column
print(df["normalized-losses"].unique())
```

```
[nan '164' '158' '192' '188' '121' '98' '81' '118' '148' '110' '145' '137'
 '101' '78' '106' '85' '107' '104' '113' '150' '129' '115' '93' '142'
 '161' '153' '125' '128' '122' '103' '168' '108' '194' '231' '119' '154'
 '74' '186' '83' '102' '89' '87' '77' '91' '134' '65' '197' '90' '94'
 '256' '95']
```

```
In [45]: # Change datatype of column
df["normalized-losses"] = df["normalized-losses"].astype("float")
```

```
In [46]: # replace missing values of column normalized-losses bu mean value
df["normalized-losses"].replace(np.nan, df["normalized-losses"].mean(), inplace
=True)
```


In [66]: `df.tail(20)`

Out[66]:

	Symbols	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	v
185	2	94	volkswagen	gas	std	four	sedan	fwd	front	
186	2	94	volkswagen	gas	std	four	sedan	fwd	front	
187	2	94	volkswagen	diesel	turbo	four	sedan	fwd	front	
188	2	94	volkswagen	gas	std	four	sedan	fwd	front	
189	3	122	volkswagen	gas	std	two	convertible	fwd	front	
190	3	256	volkswagen	gas	std	two	hatchback	fwd	front	
191	0	122	volkswagen	gas	std	four	sedan	fwd	front	
192	0	122	volkswagen	diesel	turbo	four	sedan	fwd	front	
193	0	122	volkswagen	gas	std	four	wagon	fwd	front	
194	-2	103	volvo	gas	std	four	sedan	rwd	front	
195	-1	74	volvo	gas	std	four	wagon	rwd	front	
196	-2	103	volvo	gas	std	four	sedan	rwd	front	
197	-1	74	volvo	gas	std	four	wagon	rwd	front	
198	-2	103	volvo	gas	turbo	four	sedan	rwd	front	
199	-1	74	volvo	gas	turbo	four	wagon	rwd	front	
200	-1	95	volvo	gas	std	four	sedan	rwd	front	
201	-1	95	volvo	gas	turbo	four	sedan	rwd	front	
202	-1	95	volvo	gas	std	four	sedan	rwd	front	
203	-1	95	volvo	diesel	turbo	four	sedan	rwd	front	
204	-1	95	volvo	gas	turbo	four	sedan	rwd	front	

20 rows × 26 columns



In []: *# Check the unique values, data-type of the columns with missing value
replace missing values of columns - bore, stroke, horsepower, peak-rpm by mean value and that of number of doors by the value with highest frequency*

In []:

In [53]: `#df["bore"].replace(np.nan, df["bore"].mean(), inplace=True)`

In [54]: `df['bore'] = df['bore'].astype("float")`

In [55]: `df["bore"].replace(np.nan, df["bore"].mean(), inplace=True)`

```
In [56]: df['stroke'] = df['stroke'].astype("float")
df['horsepower'] = df['horsepower'].astype("float")
df['peak-rpm'] = df['peak-rpm'].astype("float")
```

```
In [57]: df["stroke"].replace(np.nan, df["stroke"].mean(),inplace=True)
df["horsepower"].replace(np.nan, df["horsepower"].mean(),inplace=True)
df["peak-rpm"].replace(np.nan, df["peak-rpm"].mean(),inplace=True)
```

```
In [58]: print(df["num-of-doors"].value_counts())
df["num-of-doors"].value_counts().idxmax()
```

```
four      113
two        86
Name: num-of-doors, dtype: int64
```

Out[58]: 'four'

```
In [59]: df["num-of-doors"].replace(np.nan, df["num-of-doors"].value_counts().idxmax(),
inplace=True)
```

```
In [60]: df.dtypes
```

```
Out[60]: Symbols      int64
normalized-losses    float64
make                 object
fuel-type            object
aspiration           object
num-of-doors         object
body-style           object
drive-wheels         object
engine-location      object
wheel-base          float64
length              float64
width               float64
height              float64
curb-weight          int64
engine-type          object
num-of-cylinders     object
engine-size          int64
fuel-system          object
bore                 float64
stroke              float64
compression-ratio    float64
horsepower           float64
peak-rpm            float64
city-mpg             int64
highway-mpg          int64
price                object
dtype: object
```

```
In [62]: df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
df[["price"]] = df[["price"]].astype("float")
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
```

```
In [61]: df['normalized-losses'].unique()
```

```
Out[61]: array([122., 164., 158., 192., 188., 121., 98., 81., 118., 148., 110.,
145., 137., 101., 78., 106., 85., 107., 104., 113., 150., 129.,
115., 93., 142., 161., 153., 125., 128., 103., 168., 108., 194.,
231., 119., 154., 74., 186., 83., 102., 89., 87., 77., 91.,
134., 65., 197., 90., 94., 256., 95.])
```

```
In [63]: df['fuel-type'].unique()
```

```
Out[63]: array(['gas', 'diesel'], dtype=object)
```

```
In [ ]: #df['fuel-type'].replace(['gas', 'diesel'], [0,1], inplace=True)
```

```
In [67]: dummy_variable_1 = pd.get_dummies(df["fuel-type"])
dummy_variable_1.head()
```

```
Out[67]:
```

	diesel	gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

```
In [69]: dummy_variable_1.rename(columns={'gas':'fuel-type-diesel', 'diesel':'fuel-type
-gas'}, inplace=True)
dummy_variable_1.head()
```

```
Out[69]:
```

	fuel-type-diesel	fuel-type-diesel
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

```
In [70]: df = pd.concat([df, dummy_variable_1], axis=0)
```

In [71]: `df.head()`

Out[71]:

	Symbols	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4

5 rows × 28 columns



In [72]: `df.drop("fuel-type", axis = 1, inplace=True)`

In [73]: `df.head()`

Out[73]:

	Symbols	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	168.8
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	168.8
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	171.2
3	2	164	audi	std	four	sedan	fwd	front	99.8	176.6
4	2	164	audi	std	four	sedan	4wd	front	99.4	176.6

5 rows × 27 columns



In [74]: `df["Temp"] = (df["Symbols"] * 6) / 3`

In [75]: `df.head()`

Out[75]:

	Symbols	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	168.8
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	168.8
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	171.2
3	2	164	audi	std	four	sedan	fwd	front	99.8	176.6
4	2	164	audi	std	four	sedan	4wd	front	99.4	176.6

5 rows × 28 columns



In [76]: `from sklearn.linear_model import LinearRegression`

In []:

In [77]: `df.dtypes`

Out[77]:

Symbols	int64
normalized-losses	int32
make	object
aspiration	object
num-of-doors	object
body-style	object
drive-wheels	object
engine-location	object
wheel-base	float64
length	float64
width	float64
height	float64
curb-weight	int64
engine-type	object
num-of-cylinders	object
engine-size	int64
fuel-system	object
bore	float64
stroke	float64
compression-ratio	float64
horsepower	float64
peak-rpm	float64
city-mpg	int64
highway-mpg	int64
price	float64
fuel-type-diesel	uint8
fuel-type-diesel	uint8
Temp	float64
dtype:	object

```
In [79]: x = df[['horsepower', 'stroke', 'engine-size', 'highway-mpg']]  
x.head()
```

Out[79]:

	horsepower	stroke	engine-size	highway-mpg
0	111.0	2.68	130	27
1	111.0	2.68	130	27
2	154.0	3.47	152	26
3	102.0	3.40	109	30
4	115.0	3.40	136	22

```
In [80]: Linear_model = LinearRegression()
```

```
In [81]: Linear_model.fit(x, df['price'])
```

Out[81]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

```
In [82]: y_predict = Linear_model.predict(x)
          print(y_predict)
          Linear_model.score(x,df['price'])
```

```
[15468.2607231 15468.2607231 18570.95731511 10703.98308272
15644.67758135 15068.16808026 15068.16808026 15068.16808026
16252.15001326 11808.02710559 11808.02710559 19066.99689582
19066.99689582 19446.00985016 27547.40351676 27547.40351676
27800.07881966 296.78065028 5946.21943761 5946.21943761
5891.14756146 6270.1605158 9331.13713236 6270.1605158
6270.1605158 6270.1605158 9331.13713236 11681.60104394
18154.60497641 3763.71640053 6496.10639197 4361.90518528
7001.45699776 7001.45699776 7001.45699776 7001.45699776
9476.39545922 9476.39545922 9476.39545922 9476.39545922
10700.5733567 10282.06109301 10459.47201215 11945.76069191
32402.87389907 32402.87389907 47344.80961429 7433.5602371
6549.19667697 6549.19667697 6549.19667697 6549.19667697
6886.67550143 6886.67550143 6886.67550143 9495.02322876
11404.36516385 11404.36516385 11404.36516385 11404.36516385
9351.00246239 11404.36516385 16174.13147533 11087.89584215
21071.40666552 21071.40666552 21071.40666552 21071.40666552
30702.50981455 30702.50981455 40987.85253778 40481.704054
18801.85596401 6144.22180335 6523.23475769 6523.23475769
9331.13713236 11269.13625439 11428.92574104 18230.85449652
18230.85449652 18230.85449652 11428.92574104 11428.92574104
11269.13625439 11269.13625439 7207.38304304 5428.10312849
7207.38304304 7207.38304304 7207.38304304 7207.38304304
7207.38304304 7207.38304304 7207.38304304 7207.38304304
11259.60760038 11259.60760038 23048.13341016 23048.13341016
22669.12045583 22985.11493062 24817.7626075 22985.11493062
13056.73075564 15260.82260321 13056.73075564 16271.52381479
14883.9701398 15260.82260321 14883.9701398 16271.52381479
13056.73075564 15260.82260321 16567.59460955 5891.14756146
9331.13713236 6270.1605158 6270.1605158 7282.45748336
11681.60104394 18230.85449652 18569.8358209 27191.87310339
27191.87310339 27191.87310339 12590.06015127 12590.06015127
13420.15685269 13420.15685269 15326.39485555 13420.15685269
15647.79762307 15647.79762307 9106.52203714 10754.36922136
10754.36922136 10351.83709683 10857.18770262 11457.51706626
11867.88891419 12508.01827954 10983.52535406 11583.8547177
11362.5383084 13266.04418822 6541.14885072 6667.48650216
6667.48650216 6793.82415361 7425.51241084 7425.51241084
7869.04135407 7869.04135407 8350.83796505 6961.12379913
6605.66483961 8248.05430841 8248.05430841 8248.05430841
8248.05430841 10443.40165819 10443.40165819 15748.22308828
15748.22308828 15748.22308828 15748.22308828 15748.22308828
15748.22308828 11181.74863532 9401.33917833 11434.42393821
11434.42393821 11434.42393821 21733.08164174 21733.08164174
21535.58509499 20270.21388555 5189.16974076 9527.14421799
5189.16974076 9527.14421799 9527.14421799 6326.50929614
10372.30916113 10356.32902197 10356.32902197 15194.50573171
6831.85990193 10024.65510038 15956.39746875 15956.39746875
15956.39746875 15956.39746875 17218.48189583 17218.48189583
15956.39746875 18152.37865318 21961.00642401 15796.32962847
16335.41042309]
```

Out[82]: 0.7975597736149305

In [83]: `from sklearn.metrics import mean_squared_error`


```
In [84]: mse = mean_squared_error(df['price'],y_predict)
mse
```

Out[84]: 12721678.898490274

```
In [ ]:
```